

Cryoocyte Test Project

Ming Chen

10/5/2018

Import Libraries

```
library(ggplot2)    # data visualization
library(tidyverse)  # a set of module for data manipulation
library(leaps)
library(gridExtra)
```

Import Data

```
my_data = read_csv('data/train.csv')
my_test_data = read_csv('data/test.csv')
```

Linear Model selection

The data is being randomly split into training and testing sets in the 70/30 ratio. The forward selection method is implemented to determine the best final model. In the initial round of machine learning analysis, the full (with all 100 variables included) linear model has an R2 score up to **0.999**. In this section, I aim to find a simpler linear model which has an R2 score larger than **0.995** but use less features.

Forward selection

```
n = 20 # 20x cross validation: train/test = 70/30

#-----first run-----
index = sample(1:nrow(my_data), size = round(0.7*nrow(my_data)))
train = my_data[index, ]
test = my_data[-index, ]
regfit.fwd = regsubsets(y~., data = train, nvmax = 1000, method = "forward")
regfit.fwd.summary = summary(regfit.fwd)

# save train and test errors
best_num_of_variables = head(which(regfit.fwd.summary$adjr2 > 0.995), 1)
best_variables = paste0(names(coef(regfit.fwd, best_num_of_variables))[-1], collapse=',')
best_models = data.frame(best_num_of_variables, best_variables, stringsAsFactors = FALSE)
selected_columns = c(names(coef(regfit.fwd, best_num_of_variables))[-1]), 'y')
lm.fit = lm(y~., data = train[selected_columns])
train_pred = predict(lm.fit, newdata = train[selected_columns])
test_pred = predict(lm.fit, newdata = test[selected_columns])
train_error = (train$y - train_pred)^2
test_error = (test$y - test_pred)^2
errors_df = data.frame(run_01=c(train_error, test_error))
```

```

# save performance metrics
RSS_df = data.frame(run_01 = regfit.fwd.summary$rss)
RSq_df = data.frame(run_01 = regfit.fwd.summary$adjr2)
Cp_df = data.frame(run_01 = regfit.fwd.summary$cp)
BIC_df = data.frame(run_01 = regfit.fwd.summary$bic)

#-----second to n runs-----
for (i in 2:n) {
  index = sample(1:nrow(my_data), size = round(0.7*nrow(my_data)))
  train = my_data[index, ]
  test = my_data[-index, ]
  regfit.fwd = regsubsets(y~., data = train, nvmax = 1000, method = "forward")
  regfit.fwd.summary = summary(regfit.fwd)
  if (i < 10) {
    col_name = paste0('run_0', i)
  } else {
    col_name = paste0('run_', i)
  }

  # save train and test errors
  best_num_of_variables = head(which(regfit.fwd.summary$adjr2 > 0.995), 1)
  best_variables = paste0(names(coef(regfit.fwd, best_num_of_variables))[-1], collapse=',')
  best_models[i, ] = c(best_num_of_variables, best_variables)
  selected_columns = c(names(coef(regfit.fwd, best_num_of_variables)[-1]), 'y')
  lm.fit = lm(y~., data = train[selected_columns])
  train_pred = predict(lm.fit, newdata = train[selected_columns])
  test_pred = predict(lm.fit, newdata = test[selected_columns])
  train_error = (train$y - train_pred)^2
  test_error = (test$y - test_pred)^2
  errors_df[col_name] = data.frame(run_1=c(train_error, test_error))

  # save performance metrics
  RSS_df[col_name] = regfit.fwd.summary$rss
  RSq_df[col_name] = regfit.fwd.summary$adjr2
  Cp_df[col_name] = regfit.fwd.summary$cp
  BIC_df[col_name] = regfit.fwd.summary$bic
}

```

Determine which variables should be kept in the final linear model

After randomly splitting the data and implementing the forward selection 20 times, we can determine what's the best number of variables to keep and which variables they are if we want to have an R2 score up to 0.995.

The best of number of features to keep

```

# frequency of best number of variables
table(best_models$best_num_of_variables)


```

```

##
## 24 27 28 32

```

```
## 1 7 11 1
barplot(table(best_models$best_num_of_variables))
```



The most important features

```
# most important variables
highest_best_num = as.integer(names(which.max(table(best_models$best_num_of_variables))))
most_important_features = sort(table(strsplit(paste0(best_models$best_variables, collapse=','), ',')[1:
  decreasing = TRUE])[1:highest_best_num])
most_important_features
```

```
##
## x1 x15 x19 x2 x25 x29 x31 x42 x5 x59 x62 x72 x74 x76 x84 x87 x98 x99
## 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20
## x7 x53 x55 x69 x28 x3 x14 x78 x79 x80
## 17 16 14 14 13 13 12 12 11 11
```

The relationship between the number of variables and the metrics of model performance

```
# RSS
gather_RSS_df = gather(RSS_df)
gather_RSS_df['x'] = rep(1:nrow(RSS_df), n)
p1 = ggplot(data = gather_RSS_df, aes(x=x, y=value, color=key)) +
  geom_line() +
  geom_vline(xintercept = highest_best_num) +
  xlab("Number of Variables") +
  ylab("RSS") +
  theme(legend.position="none")
```

```

# RSq
gather_RSq_df = gather(RSq_df)
gather_RSq_df['x'] = rep(1:nrow(RSq_df), n)
p2 = ggplot(data = gather_RSq_df, aes(x=x, y=value, color=key)) +
  geom_line() +
  geom_vline(xintercept = highest_best_num) +
  xlab("Number of Variables") +
  ylab("Adjusted R Squared") +
  theme(legend.position="none")

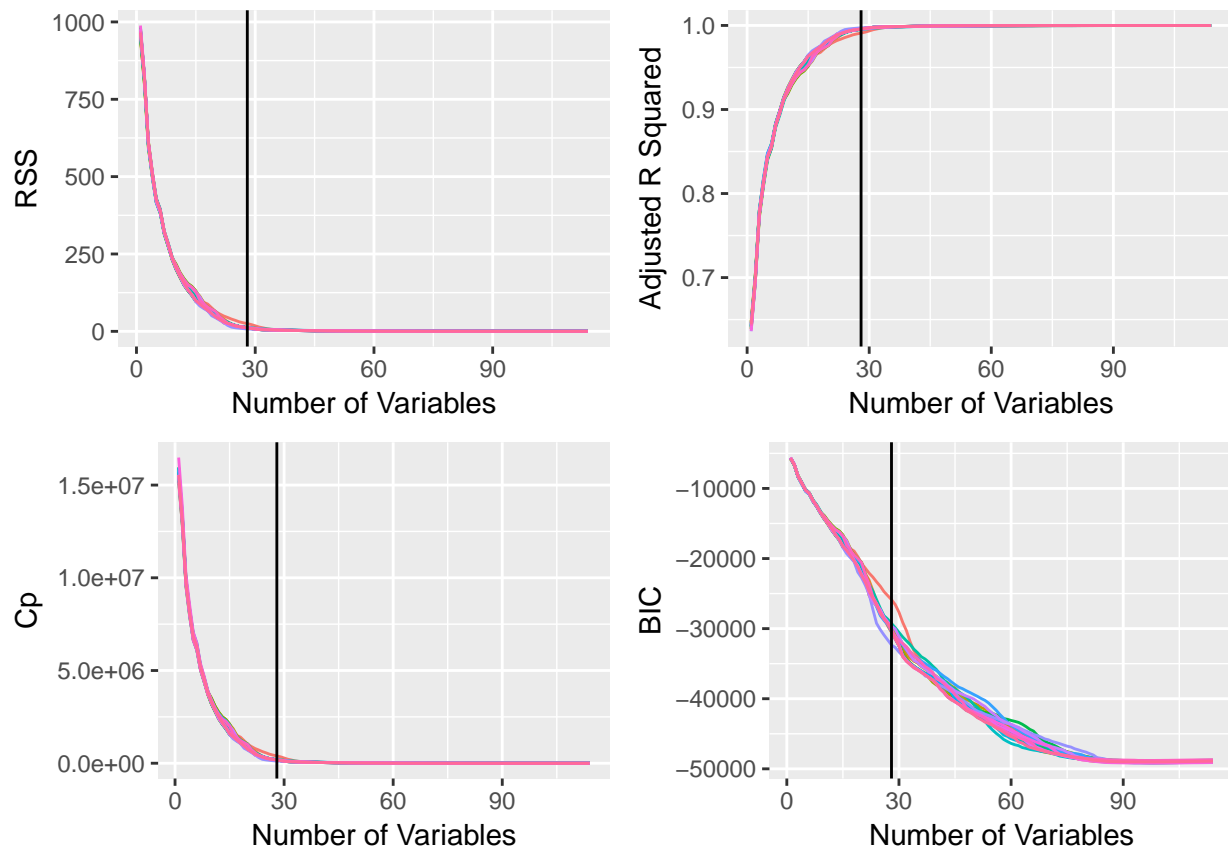
# grid.arrange(p1, p2, ncol=2)

# Cp
gather_Cp_df = gather(Cp_df)
gather_Cp_df['x'] = rep(1:nrow(Cp_df), n)
p3 = ggplot(data = gather_Cp_df, aes(x=x, y=value, color=key)) +
  geom_line() +
  geom_vline(xintercept = highest_best_num) +
  xlab("Number of Variables") +
  ylab("Cp") +
  theme(legend.position="none")

# BIC
gather_BIC_df = gather(BIC_df)
gather_BIC_df['x'] = rep(1:nrow(BIC_df), n)
p4 = ggplot(data = gather_BIC_df, aes(x=x, y=value, color=key)) +
  geom_line() +
  geom_vline(xintercept = highest_best_num) +
  xlab("Number of Variables") +
  ylab("BIC") +
  theme(legend.position="none")

grid.arrange(p1, p2, p3, p4, nrow=2, ncol=2)

```




Training vs. Test errors

The plot below shows that the training and test errors are very low and pretty close to each other across all 20 runs, which means that model fits both the training and test data sets well.

```
# Erros DataFrame
errors_df$error_type = c(rep('train', nrow(train)), rep('test', nrow(test)))
gather_errors_df = gather(errors_df, key = "key", value = "value",
                           setdiff(names(errors_df), 'error_type'))

ggplot(data = gather_errors_df, aes(x = key, y = value, color=error_type)) +
  geom_boxplot() +
  xlab(paste0(n, ' runs')) +
  ylab('MSE') +
  scale_color_discrete(name = "Error Type") +
  theme(axis.text.x=element_blank(),
        axis.ticks.x=element_blank())
```



Final Linear Model

```
final_lm = lm(y~., my_data[c(names(most_important_features), 'y')])
# remove non significant terms
update_final_lm = update(final_lm, . ~ . -1)
final_lm_summary = summary(update_final_lm)
final_lm_summary
```

```
##
## Call:
## lm(formula = y ~ x1 + x15 + x19 + x2 + x25 + x29 + x31 + x42 +
##      x5 + x59 + x62 + x72 + x74 + x76 + x84 + x87 + x98 + x99 +
##      x7 + x53 + x55 + x69 + x28 + x3 + x14 + x78 + x79 + x80 -
##      1, data = my_data[c(names(most_important_features), "y")])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.163727 -0.030440 -0.000535  0.031158  0.153906
##
## Coefficients:
##      Estimate Std. Error t value Pr(>|t|)
## x1    60.4968     0.4121  146.78  <2e-16 ***
## x15   15.9961     0.4926   32.47  <2e-16 ***
## x19  108.4670     0.6177  175.60  <2e-16 ***
## x2    93.0995     0.5141  181.08  <2e-16 ***
```

```

## x25 49.2793      0.5781   85.25   <2e-16 ***
## x29 122.7685     0.7440  165.00   <2e-16 ***
## x31 68.2603      0.6282  108.66   <2e-16 ***
## x42 95.4224      0.6816  139.99   <2e-16 ***
## x5  91.1176      0.6440  141.50   <2e-16 ***
## x59 68.9285      0.5926  116.32   <2e-16 ***
## x62 11.9256      0.5778   20.64   <2e-16 ***
## x72 84.4303      0.6670  126.59   <2e-16 ***
## x74 77.5597      0.5447  142.39   <2e-16 ***
## x76 39.7352      0.8160   48.69   <2e-16 ***
## x84 69.4174      0.6315  109.93   <2e-16 ***
## x87 45.4966      0.6790   67.00   <2e-16 ***
## x98 94.4588      0.9910   95.31   <2e-16 ***
## x99 50.1494      0.7275   68.94   <2e-16 ***
## x7  22.0609      0.5439   40.56   <2e-16 ***
## x53 18.7576      0.4198   44.68   <2e-16 ***
## x55 -27.7065     0.6633  -41.77   <2e-16 ***
## x69 -31.9632     0.4915  -65.03   <2e-16 ***
## x28 76.8208      0.6607  116.27   <2e-16 ***
## x3  25.7173      0.6238   41.23   <2e-16 ***
## x14 28.8451      0.5550   51.98   <2e-16 ***
## x78 27.0942      0.6231   43.48   <2e-16 ***
## x79 45.1815      0.7092   63.71   <2e-16 ***
## x80 21.9096      0.5603   39.10   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04514 on 7949 degrees of freedom
## (23 observations deleted due to missingness)
## Multiple R-squared:  0.9958, Adjusted R-squared:  0.9958
## F-statistic: 6.802e+04 on 28 and 7949 DF,  p-value: < 2.2e-16

```

The final linear model includes 28 features: x1, x15, x19, x2, x25, x29, x31, x42, x5, x59, x62, x72, x74, x76, x84, x87, x98, x99, x7, x53, x55, x69, x28, x3, x14, x78, x79, x80.